

A simple example

- ✓ Inventory ordering problem
- ✓ Data:
 - ✓ Starting stock level
 - ✓ Maximum stock level
 - ✓ Demand in each time period (deterministic)
 - ✓ Total cost of ordering n units of stock = $F(n)$, for some given function F
 - ✓ Typically non-linear, increasing
 - ✓ Cost per unit per time period of holding stock
 - ✓ Stock out cost
- ✓ Determine:
 - ✓ Quantity of stock to order in each period

Let's get our notation straight

S Discrete state space

A Discrete action space

Stages indexed by t : This often corresponds to time periods

Transition function: $S_{t+1} = S^M(S_t, a_t)$

The state we are in at stage $t+1$ depends on the previous state and action

In our example, the new inventory level depends on the previous level, quantity ordered and the demand

Contribution function: $C_t(S_t, a_t)$

The immediate cost (or value) of making decision a_t in state S_t . F in our example.

Value function: $V_t(S_t) = \min_{a_t} \{C_t(S_t, a_t) + V_{t+1}(S_{t+1})\}$

Total cost (or value) of making all the best decisions from here to the end

Solving simple dynamic programming problems is easy

- ✓ Code recursive function directly
- ✓ “Memo-ise” values for efficiency
- ✓ Return value function and the argument that achieves the optimal value

Extending to stochastic problems is sometimes easy

✓ Our example:

- ✓ Demand is no longer constant, but rather given by a known, discrete probability distribution

✓ Notation:

- ✓ Transition matrix: $p_t(S_{t+1}|S_t, a_t)$
Probability that if we are in state S_t and take action a_t that we will next be in state S_{t+1}
For our example this can be calculated from the demand distribution and our action

- ✓ Value function:
$$V_t(S_t) = \min_{a_t} \left\{ C_t(S_t, a_t) + \sum_{s' \in S} p_t(s'|S_t, a_t) V_{t+1}(s') \right\}$$

alternatively:
$$V_t(S_t) = \min_{a_t} \left\{ C_t(S_t, a_t) + E[V_{t+1}(S_{t+1})|S_t, a_t] \right\}$$

Solving simple stochastic dynamic programming problems is easy

- ✓ Code recursive function directly
- ✓ “Memo-ise” values for efficiency
- ✓ Return value function and the argument that achieves the optimal value
- ✓ The answer is an expected value and a **policy**. For each possible state, it specifies the optimal action.

Some problems are too hard to solve exactly

- ✓ Consider our example expanded to 10 product types:
 - ✓ If we have 1000 maximum units in stock for each product, we have 1001^{10} possible states
 - ✓ If demand for each can range from 0 to 200, we have 201^{10} possible outcomes
 - ✓ If we can order between up to 500 units at a time, we have 501^{10} possible actions
- ✓ These are the three curses of dimensionality:
 - ✓ State space – traditional curse of dimensionality for deterministic problems
 - ✓ Outcome space – we may not be able to compute our expectation
 - ✓ Decision space – LP and MIP regularly handle very large decision spaces